

sklearn-rcd180001

November 4, 2022

1 ML with SKLearn

Joshua Durana rcd180001

```
[ ]: import pandas as pd
import sklearn as sk
import seaborn as sb
import numpy as np
```

1.1 Load Data

```
[ ]: autodf = pd.read_csv('Data/Auto.csv')
print("Head\n", autodf.head())
print("Dimensions: ", autodf.shape)
```

Head

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70.0	
1	15.0	8	350.0	165	3693	11.5	70.0	
2	18.0	8	318.0	150	3436	11.0	70.0	
3	16.0	8	304.0	150	3433	12.0	70.0	
4	17.0	8	302.0	140	3449	NaN	70.0	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

Dimensions: (392, 9)

1.2 Data Exploration

```
[ ]: autodf.describe()
```

```
[ ]:
count      mpg  cylinders  displacement  horsepower      weight  \
count  392.000000  392.000000   392.000000  392.000000  392.000000
mean    23.445918    5.471939   194.411990  104.469388  2977.584184
```

std	7.805007	1.705783	104.644004	38.491160	849.402560
min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.000000	4.000000	105.000000	75.000000	2225.250000
50%	22.750000	4.000000	151.000000	93.500000	2803.500000
75%	29.000000	8.000000	275.750000	126.000000	3614.750000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	acceleration	year	origin
count	391.000000	390.000000	392.000000
mean	15.554220	76.010256	1.576531
std	2.750548	3.668093	0.805518
min	8.000000	70.000000	1.000000
25%	13.800000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.050000	79.000000	2.000000
max	24.800000	82.000000	3.000000

- The mpg's range is 37 and the mean is 23.
- The cylinder's range is 4 and the mean is 5.5 cylinders.
- The displacement's range is 387 and the mean is 194.4
- The horsepower's range is 184 and the mean is 104.5
- The weight's range is 3527 and the mean is 2977.6
- The acceleration range is 16.8 and the mean is 15.6
- The year's range is 12 and the mean is 76
- The origin's range is 2 and the mean is 1.6

1.3 Data Types

```
[ ]: print("Old\n",autodf.dtypes)

#Change Columns to Categorical
autodf.cylinders = autodf.cylinders.astype('category')
autodf.cylinders = autodf.cylinders.cat.codes

autodf.origin = pd.Categorical(autodf.origin)

print("\nNew\n", autodf.dtypes)
```

```
Old
mpg                float64
cylinders           int64
displacement       float64
horsepower         int64
weight             int64
acceleration       float64
year              float64
origin            int64
```

```

name          object
dtype: object

New
  mpg          float64
  cylinders      int8
  displacement  float64
  horsepower    int64
  weight        int64
  acceleration  float64
  year          float64
  origin        category
  name          object
dtype: object

```

1.4 NA Values

```

[ ]: #Count NA values
print(autodf.isnull().sum())

#Drop NA rows
autodf = autodf.dropna()

print("\n", autodf.isnull().sum())
print("Dimensions: ", autodf.shape)

```

```

mpg          0
cylinders    0
displacement 0
horsepower   0
weight       0
acceleration 1
year         2
origin       0
name         0
dtype: int64

```

```

  mpg          0
  cylinders    0
  displacement 0
  horsepower   0
  weight       0
  acceleration 0
  year         0
  origin       0
  name         0
dtype: int64
Dimensions: (389, 9)

```

1.5 Modify Columns

```
[ ]: #Create mpg high column
autodf['mpg_high'] = [1 if m > 23.445918 else 0 for m in autodf['mpg']]

#Drop mpg and name columns
autodf = autodf.drop(columns=['mpg', 'name'])
autodf.head()
```


```
[ ]:   cylinders  displacement  horsepower  weight  acceleration  year  origin  \
0         4         307.0         130    3504         12.0    70.0      1
1         4         350.0         165    3693         11.5    70.0      1
2         4         318.0         150    3436         11.0    70.0      1
3         4         304.0         150    3433         12.0    70.0      1
6         4         454.0         220    4354          9.0    70.0      1

      mpg_high
0           0
1           0
2           0
3           0
6           0
```

1.6 Graphical Data Exploration

```
[ ]: #MPG High Catplot
sb.catplot(x = "mpg_high", kind = "count", data=autodf)
```


```
[ ]: <seaborn.axisgrid.FacetGrid at 0x7f6df2508df0>
```



It seems that mpg high seems evenly distributed, but there's more low_mpg cars

```
[ ]: #Relplot  
sb.relplot(x = "horsepower", y = "weight", hue = "mpg_high", data = autodef)
```


```
[ ]: <seaborn.axisgrid.FacetGrid at 0x7f6df265ce50>
```



The mpg_high seems clustered. High mpg is clustered on high horsepower and weight, while low mpg is clustered on low horsepower and weight. This might show that decision trees might be a good model to use.

```
[ ]: #Boxplot  
sb.boxplot(x = "mpg_high", y = "weight", data = autodf)
```

```
[ ]: <AxesSubplot: xlabel='mpg_high', ylabel='weight'>
```



The low mpg cars seems to have a higher average weight than high mpg cars. The quartiles and range seems to be larger on the low mpg cars compared to the high mpg cars. Mpg high cars also seem to have a tiny bit of outliers. Most likely, the higher the weight of the car the more likely the car is mpg_low.

1.7 Split Data to Train and Test

```
[ ]: from sklearn.model_selection import train_test_split

#Obtain predictors and target columns
predictors = autodf.drop(columns=['mpg_high'])
target = autodf.mpg_high

#Split Data
predictorTrain, predictorTest, targetTrain, targetTest = \
    ↪train_test_split(predictors, target, test_size=.2, random_state=1234)

#Output Dimensions
print("Train Dimensions:", predictorTrain.shape)
print("Test Dimensions:", predictorTest.shape)
```

Train Dimensions: (311, 7)
Test Dimensions: (78, 7)

1.8 Logistic Regression

```
[ ]: from sklearn.linear_model import LogisticRegression

#Make model
lr = LogisticRegression()
lr.fit(predictorTrain, targetTrain)
print("Score: ", lr.score(predictorTrain, targetTrain))
```

Score: 0.9067524115755627

/home/pretaxend/.local/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
[ ]: from sklearn.metrics import classification_report

#Predict
predictions = lr.predict(predictorTest)

#Metrics
print(classification_report(targetTest, predictions))
```

	precision	recall	f1-score	support
0	0.98	0.80	0.88	50
1	0.73	0.96	0.83	28
accuracy			0.86	78
macro avg	0.85	0.88	0.85	78
weighted avg	0.89	0.86	0.86	78

1.9 Decision Tree

```
[ ]: from sklearn.tree import DecisionTreeClassifier
from sklearn import tree

#Make Model
dt = DecisionTreeClassifier()
dt.fit(predictorTrain, targetTrain)
print("Score: ", dt.score(predictorTrain, targetTrain))
```

Score: 1.0

```
[ ]: #Predict
predictions = dt.predict(predictorTest)

#Metrics
print(classification_report(targetTest, predictions))

#Plot Tree
tree.plot_tree(dt)
```

	precision	recall	f1-score	support
0	0.94	0.88	0.91	50
1	0.81	0.89	0.85	28
accuracy			0.88	78
macro avg	0.87	0.89	0.88	78
weighted avg	0.89	0.88	0.89	78

```
[ ]: [Text(0.6666666666666666, 0.9444444444444444, 'X[0] <= 2.5\ngini = 0.5\nsamples = 311\nvalue = [153, 158]'),
Text(0.4583333333333333, 0.8333333333333334, 'X[2] <= 101.0\ngini = 0.239\nsamples = 173\nvalue = [24, 149]'),
Text(0.3055555555555556, 0.7222222222222222, 'X[5] <= 75.5\ngini = 0.179\nsamples = 161\nvalue = [16, 145]'),
Text(0.1666666666666666, 0.6111111111111112, 'X[1] <= 119.5\ngini = 0.362\nsamples = 59\nvalue = [14, 45]'),
Text(0.0555555555555556, 0.5, 'X[4] <= 13.75\ngini = 0.159\nsamples = 46\nvalue = [4, 42]'),
Text(0.02777777777777778, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(0.0833333333333333, 0.3888888888888889, 'X[3] <= 2683.0\ngini = 0.087\nsamples = 44\nvalue = [2, 42]'),
Text(0.0555555555555556, 0.2777777777777778, 'X[3] <= 2377.0\ngini = 0.045\nsamples = 43\nvalue = [1, 42]'),
Text(0.02777777777777778, 0.1666666666666666, 'gini = 0.0\nsamples =
```

```


38\nvalue = [0, 38]'),
Text(0.08333333333333333, 0.16666666666666666, 'X[3] <= 2385.0\ngini =
0.32\nsamples = 5\nvalue = [1, 4]'),
Text(0.05555555555555555, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue
= [1, 0]'),
Text(0.11111111111111111, 0.05555555555555555, 'gini = 0.0\nsamples = 4\nvalue =
[0, 4]'),
Text(0.11111111111111111, 0.27777777777777778, 'gini = 0.0\nsamples = 1\nvalue =
[1, 0]'),
Text(0.27777777777777778, 0.5, 'X[3] <= 2567.0\ngini = 0.355\nsamples =
13\nvalue = [10, 3]'),
Text(0.25, 0.38888888888888889, 'X[3] <= 2429.5\ngini = 0.469\nsamples =
8\nvalue = [5, 3]'),
Text(0.22222222222222222, 0.27777777777777778, 'X[4] <= 15.75\ngini =
0.278\nsamples = 6\nvalue = [5, 1]'),
Text(0.19444444444444444, 0.16666666666666666, 'X[2] <= 88.0\ngini =
0.5\nsamples = 2\nvalue = [1, 1]'),
Text(0.16666666666666666, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue
= [1, 0]'),
Text(0.22222222222222222, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.25, 0.16666666666666666, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(0.27777777777777778, 0.27777777777777778, 'gini = 0.0\nsamples = 2\nvalue =
[0, 2]'),
Text(0.30555555555555556, 0.38888888888888889, 'gini = 0.0\nsamples = 5\nvalue =
[5, 0]'),
Text(0.44444444444444444, 0.61111111111111112, 'X[3] <= 3250.0\ngini =
0.038\nsamples = 102\nvalue = [2, 100]'),
Text(0.38888888888888889, 0.5, 'X[3] <= 2880.0\ngini = 0.02\nsamples =
100\nvalue = [1, 99]'),
Text(0.36111111111111111, 0.38888888888888889, 'gini = 0.0\nsamples = 94\nvalue =
[0, 94]'),
Text(0.41666666666666667, 0.38888888888888889, 'X[3] <= 2920.0\ngini =
0.278\nsamples = 6\nvalue = [1, 5]'),
Text(0.38888888888888889, 0.27777777777777778, 'gini = 0.0\nsamples = 1\nvalue =
[1, 0]'),
Text(0.44444444444444444, 0.27777777777777778, 'gini = 0.0\nsamples = 5\nvalue =
[0, 5]'),
Text(0.5, 0.5, 'X[1] <= 151.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(0.47222222222222222, 0.38888888888888889, 'gini = 0.0\nsamples = 1\nvalue =
[1, 0]'),
Text(0.52777777777777778, 0.38888888888888889, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.61111111111111112, 0.72222222222222222, 'X[4] <= 14.45\ngini =
0.444\nsamples = 12\nvalue = [8, 4]'),
Text(0.58333333333333334, 0.61111111111111112, 'X[5] <= 76.0\ngini =
0.444\nsamples = 6\nvalue = [2, 4]'),

```

```

Text(0.5555555555555556, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(0.6111111111111112, 0.5, 'X[1] <= 138.5\ngini = 0.444\nsamples = 3\nvalue
= [2, 1]'),
Text(0.5833333333333334, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue =
[2, 0]'),
Text(0.6388888888888888, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.6388888888888888, 0.6111111111111112, 'gini = 0.0\nsamples = 6\nvalue =
[6, 0]'),
Text(0.875, 0.8333333333333334, 'X[5] <= 79.5\ngini = 0.122\nsamples =
138\nvalue = [129, 9]'),
Text(0.8055555555555556, 0.7222222222222222, 'X[4] <= 21.6\ngini =
0.045\nsamples = 129\nvalue = [126, 3]'),
Text(0.7777777777777778, 0.6111111111111112, 'X[3] <= 2737.0\ngini =
0.031\nsamples = 128\nvalue = [126, 2]'),
Text(0.7222222222222222, 0.5, 'X[5] <= 76.0\ngini = 0.444\nsamples = 3\nvalue =
[2, 1]'),
Text(0.6944444444444444, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue =
[2, 0]'),
Text(0.75, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.8333333333333334, 0.5, 'X[2] <= 83.0\ngini = 0.016\nsamples = 125\nvalue
= [124, 1]'),
Text(0.8055555555555556, 0.3888888888888889, 'X[1] <= 225.0\ngini =
0.375\nsamples = 4\nvalue = [3, 1]'),
Text(0.7777777777777778, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.8333333333333334, 0.2777777777777778, 'gini = 0.0\nsamples = 3\nvalue =
[3, 0]'),
Text(0.8611111111111112, 0.3888888888888889, 'gini = 0.0\nsamples = 121\nvalue
= [121, 0]'),
Text(0.8333333333333334, 0.6111111111111112, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.9444444444444444, 0.7222222222222222, 'X[6] <= 1.5\ngini =
0.444\nsamples = 9\nvalue = [3, 6]'),
Text(0.9166666666666666, 0.6111111111111112, 'X[1] <= 247.0\ngini =
0.48\nsamples = 5\nvalue = [3, 2]'),
Text(0.8888888888888888, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(0.9444444444444444, 0.5, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(0.9722222222222222, 0.6111111111111112, 'gini = 0.0\nsamples = 4\nvalue =
[0, 4]')

```



1.10 Neural Network

```
[ ]: from sklearn import preprocessing
from sklearn.neural_network import MLPClassifier

#Normalize Data
scale = preprocessing.StandardScaler().fit(predictorTrain)
predictorTrainScaled = scale.transform(predictorTrain)
predictorTestScaled = scale.transform(predictorTest)

#Train
nnlbfgs = MLPClassifier(solver = 'lbfgs', hidden_layer_sizes = (3,2), max_iter=
    500, random_state = 1234)
nnlbfgs.fit(predictorTrainScaled, targetTrain)
print("Score: ", nnlbfgs.score(predictorTrainScaled, targetTrain))
```

Score: 0.9517684887459807

```
[ ]: #Predictions
predictions = nnlbfgs.predict(predictorTestScaled)

#Metrics
print(classification_report(targetTest, predictions))
```

	precision	recall	f1-score	support
0	0.92	0.88	0.90	50
1	0.80	0.86	0.83	28
accuracy			0.87	78
macro avg	0.86	0.87	0.86	78
weighted avg	0.87	0.87	0.87	78

```
[ ]: #Different Topology
nnsgd = MLPClassifier(solver = 'sgd', hidden_layer_sizes=(3,), max_iter=1000,
    ↪random_state=1234)
nnsgd.fit(predictorTrainScaled, targetTrain)

#Score
print("Score: ", nnsgd.score(predictorTrainScaled, targetTrain))
```

Score: 0.9003215434083601

```
[ ]: #Predictions
predictions = nnsgd.predict(predictorTestScaled)

#Metrics
print(classification_report(targetTest, predictions))
```

	precision	recall	f1-score	support
0	0.93	0.80	0.86	50
1	0.71	0.89	0.79	28
accuracy			0.83	78
macro avg	0.82	0.85	0.83	78
weighted avg	0.85	0.83	0.84	78

The initial model has a higher accuracy most likely due to having more nodes. Most likely the model using the 'sgd' solver underfitted the data.

1.11 Analysis

The decision tree was the more accurate algorithm with an accuracy of 88%, most likely due to the different factors being clustered together. The highest precision for mpg_low is for logistic regression with .98, while the highest precision for mpg_high is for decision trees with .81. The best recall for mpg_low is decions trees with .88, while logistic regression has the best recall for mpg_high with .96.

The decision tree seems to be the better performing algorithm, most likely due to the 2 different factors being clustered together. The neural networks most likely underperformed due to not

having as much data to work on. The logistic regression model seemed to perform well due to some predictors having linear relationships.

I personally prefer working in Python than R with ML. While R has better functionality regarding dataframes, the Pandas library still performs well and is easy to use. Sklearn seems to be much more user friendly for creating ML models than R.